

H1 – Testing standard benchmark functions Hill-Climbing and Simulated annealing

Dincă Georgian(2E2) - Brahă Petru(2E3)

October 29, 2024

1 Abstract

This work contains a comparative analysis of local search strategies applied to four optimization benchmark functions: **Sphere**, **Rastrigin**, **Schwefel** and **Michalewicz**. Three search strategies are evaluated - **Best Improvement**, **First Improvement**, and **Simulated annealing** - across the following problem dimensions (**10D**, **30D**, and **100D**), with an additional investigation of **Worst Improvement** approach for the lower-dimensional cases (**10D** and **30D**). In addition to solution quality, the execution times of each strategy were recorded and compared. Findings from these experiments are presented and discussed.

2 Motivation

The program's target is to find the smallest optimum for the functions mentioned. What could be an obstacle?

Achieving ideal outcomes using a deterministic design inflicts some disadvantages. Its performance does not hold when faced with multiple local optimum points and complex landscapes. Studying the global optima over an interval shapes the following problem: the infinite amount of real values in the interval. How is it supposed to analyze such a scenario? Of course, there exist several exponential procedures, but the overall context is not attainable. In the previous paper, we argued that this kind of algorithm (a deterministic one) is time-consuming and not feasible in practice. By contrast, the concepts of heuristics were also visited. The cues captured guided us to examine the standard functions by changing the definition of a solution: the perfect result is not necessary, and a close call to it in a much quicker way is preferred. Time costs can be notably reduced without dramatically sacrificing the quality of the output. We nominate that the quality is satisfactory in the following implementation.

Note: The precision chosen for the outputs for the findings of this experiment.

3 Method

3.1 Setup

The algorithms have been implemented to leverage GPU acceleration using NVIDIA's CUDA framework. The system used to run the tests was equipped with an RTX 4090 graphics card, mobile version. The parallel nature of local search algorithms, where multiple independent search iterations can be executed simultaneously, makes them particularly well-suited for GPU implementation due to their architecture having a huge number of cores and threads that can run in parallel. Each search iteration operates independently on its own thread, which parallelizes the workload efficiently.

Our GPU implementation utilizes a thread organization scheme of 32 threads per block, aligning with the NVIDIA warp size for optimal execution efficiency. This configuration minimizes thread divergence and maximizes memory coalescing, concluding in improved computational performance. Each experimental configuration executes 20,000 iterations per sample, with 30 independent samples for statistical robustness.

To attempt function evaluation, we employed two popular optimization algorithms: Hill Climbing and Simulated annealing (hybridization of HC best improvement); both are covered in the next subsections.

3.2 Hill Climbing implementation

The algorithm manipulates the vast search interval of the function uniquely. First of all, it calculates N , a precise cardinal of random candidate numbers inside the domain, using the formula: $N = (b - a) * 10^p$, where a and b are lower/upper bounds of the interval in question, and p being the established precision. The magic begins with the representation of these N values: as bit strings. The random number generation is, in fact, a random bit generator; we used NVIDIA's pseudo-random number generator algorithm. Constructing a bit string requires as input the maximum length, which was calculated using $\log_2(N) + 1$. Its conversion process is ironed out by the following fitness function:

$$fitness(bitstring) = a + decimal(bitstring) * (b - a) / (2^N - 1).$$

The primary goal is to identify a peak solution by continually moving towards neighboring solutions with higher fitness outcomes. A neighbor of a bit string is a copy of the current representation but with only one flipped bit. In other words, for every bit there is associated with a neighbor; this can be understood as an iteration of the string. Depending on the improvement type, a potentially more desirable neighbor is assigned to the candidate variable. For only one run of the algorithm, improvements are performed until there is no better value in the neighborhood - the end condition. The real values are then computed and the result of the benchmark functions concerning these computations is returned.

3.3 Simulated annealing implementation

Our implementation features a hybrid approach that combines traditional SA with a Best Improvement local search phase after the temperature drops and it can no longer escape the local optimum. The initial temperature T_0 is dynamically calculated based on the dimension number using the formula $T_0 = \left\lceil \frac{fmax * n}{\ln(0.8)} \right\rceil$, where n is the number of dimensions. $fmax$ is the aproximative maximum value of the function for one dimension and assures that the temperature can be high enough for exploration:

- Sphere's is very simple and doesn't need optimized temperature initialization
- Rastrigin's - 40
- Schwefel's - 840
- Michalewicz's - 2

This adaptive initial temperature ensures appropriate scaling of the acceptance probability across different problem dimensions.

The cooling schedule is implemented with a geometric decay and a cooling rate of 0.985, maintaining a balance between exploration and exploitation. The algorithm progresses until either the temperature falls below a threshold of $T_0 \times 10^{-8}$, or when the search stagnates for four consecutive temperature changes. For each temperature, we attempt to achieve $20 * n$ successful moves, where n is the number of dimensions, with a maximum of $200 * n$ total attempts per temperature to prevent excessive computation in flat regions of the search space.

The solution representation uses a binary encoding, with neighborhood moves implemented as single bit-flips. Move acceptance follows the standard Metropolis criterion:

- Improvements are always accepted
- Deteriorating moves are accepted with probability $\exp\left(-\frac{|\Delta f|}{T}\right)$, where Δf is the absolute fitness difference and T is the current temperature

3.4 Interpretation

The experimental results demonstrate that **Best Improvement** consistently outperforms **First Improvement** in solution quality across all test functions, particularly in higher dimensions, however at a higher computational cost. **Simulated annealing** exhibits superior performance by escaping local optima, especially notable for **Rastrigin** and **Schwefel** functions, yielding better outputs when compared against **Best Improvement**. Finally, **Worst Improvement** showed inferior results at a higher computational cost than all other strategies.

4 Results

4.1 Sphere

$$f(\mathbf{x}) = \sum_{i=1}^n x_i^2, x_i \in [-5.12, 5.12]$$




Figure 1: Sphere's 2-dimensional graph function [3]

D	σ	avg. time	avg. sol.	best sol.
10	0	89ms	0	0
30	0	2.36s	0	0
100	0	61.539s	0	0

Table 1: Best improvement

D	σ	avg. time	avg. sol.	best sol.
10	0	50ms	0	0
30	0	1.21s	0	0
100	0	32.58s	0	0

Table 2: First improvement

D	σ	avg. time	avg. sol.	best sol.
10	0	200ms	0	0
30	0	4.98s	0	0

Table 3: Worst improvement

D	σ	avg. time	avg. sol.	best sol.
10	0	310ms	0	0
30	0	4.97s	0	0
100	0	152.95s	0	0

Table 4: Simulated annealing

4.2 Rastrigin

$$f(x) = A \cdot n + \sum_{i=1}^n [x_i^2 - 10 \cdot \cos(2\pi x_i)], x_i \in [-5.12, 5.12]$$




Figure 2: Rastrigin's 2-dimensional graph function [4]

D	σ	avg. time	avg. sol.	best sol.
10	0.53	307ms	1.96332	0.99496
30	2.30	7.74s	22.49692	16.15900
100	5.25	271.91s	124.04872	103.27991

Table 1: Best improvement

D	σ	avg. time	avg. sol.	best sol.
10	0.79	182ms	3.42327	1.23582
30	2.63	4.37s	29.97694	22.84631
100	4.99	154.12s	155.08213	141.60620

Table 2: First improvement

D	σ	avg. time	avg. sol.	best sol.
10	0.91	2.51s	5.12165	3.00004
30	3.23	57.13s	38.18328	29.23457

Table 3: Worst improvement

D	σ	avg. time	avg. sol.	best sol.
10	0	19.81s	0	0
30	0.86	171.31s	6.20453	3.99499
100	1.94	19.16min	47.64466	40.26256

Table 4: Simulated annealing

4.3 Schwefel

$$f(\mathbf{x}) = -\sum_{i=1}^n x_i \cdot \sin\left(\sqrt{|x_i|}\right), x_i \in [-500, 500]$$




Figure 3: Schwefel's 2-dimensional graph function [5]

D	σ	avg. time	avg. sol.	best sol.
10	18.4	755ms	17.70524	0.20936
30	98.35	19.813s	894.62572	696.92139
100	230.35	7,62min	5747.60089	5326.42004

Table 1: Best improvement

D	σ	avg. time	avg. sol.	best sol.
10	45.79	434ms	120.65099	34.54979
30	111.23	11.18s	1464.46143	1186.97187
100	217.08	261.82s	7615.11165	6976.11061

Table 2: First improvement

D	σ	avg. time	avg. sol.	best sol.
10	17.35	7.59s	212.34472	161.40279
30	108.91	118.25s	1550.94308	1262.86526

Table 3: Worst improvement

D	σ	avg. time	avg. sol.	best sol.
10	0	22.98s	0.00037	0.00014
30	0.07	173.01s	0.38921	0.21216
100	41.02	25.93min	126.46816	71.91137

Table 4: Simulated annealing

4.4 Michalewicz

$$f(\mathbf{x}) = -\sum_{i=1}^n \sin(x_i) \cdot \left(\sin\left(ix_i^2/\pi\right)\right)^{2m}, m = 10, x_i \in [0, \pi]$$




Figure 4: Michalewicz's 2-dimensional graph function [6]

D	σ	avg. time	avg. sol.	best sol.
10	0.04	1.50s	-9.56506	-9.65221
30	0.28	33.42s	-27.61249	-28.56555
100	0.51	9.73min	-88.04727	-89.11685

Table 1: Best improvement

D	σ	avg. time	avg. sol.	best sol.
10	0.05	865ms	-9.50893	-9.61160
30	0.18	19.48s	-27.02483	-27.45056

Table 2: First improvement

D	σ	avg. time	avg. sol.	best sol.
10	0.12	7.39s	-9.22340	-9.42528
30	0.29	104.44s	-24.57051	-25.19599

Table 3: Worst improvement

D	σ	avg. time	avg. sol.	best sol.
10	0	24.69s	-9.65930	-9.66015
30	0.07	210.07s	-29.11762	-29.27980
100	0.17	37.3min	-95.99791	-96.32860

Table 4: Simulated annealing

5 Conclusions

Looking into our results, it can safely be stated that the Hill Climbing algorithm is efficient and succeeds in delivering significant values, in the search for the global minimum. Comparing it with the deterministic approach, the time is much less than the brute-force traversal of the graph function. The different improvement settings showcase diverse trade-offs between these designs which can be further exploited.

An exemplary observation would be the Simulated annealing contributions; they usually bring better solutions that are scalable for real competitive environments. It's useful in concrete cases, even though considerable time intervals were detected.

Overall, this study provides valuable insights into the application of optimization algorithms in solving complex optimization problems and contributes to the understanding of their performance across different dimensions.

References

- [1] Course page
Hill Climbing documentation. <https://profs.info.uaic.ro/eugen.croitoru/teaching/ga/>
- [2] Nvidia Cuda coalesced memory
<https://docs.nvidia.com/cuda/cuda-c-best-practices-guide/index.html#coalesced-access-to-global-memory>
- [3] Simon Fraser University
Sphere's function. <https://www.sfu.ca/~ssurjano/spheref.html>
- [4] Simon Fraser University
Rastrigin's function. <https://www.sfu.ca/~ssurjano/rastr.html>
- [5] Simon Fraser University
Schwefel's function. <https://www.sfu.ca/~ssurjano/schwef.html>
- [6] Simon Fraser University
Michalewicz's function. <https://www.sfu.ca/~ssurjano/michal.html>